

Amplified Oracle: Probably Better Methods for Rewarding Forecasters

Nuño Sempere*

April 15, 2022

Motivation

In [Alignment Problems With Current Forecasting Platforms](#), Sempere and Lawsen outline a variety of problems with current forecasting platforms, whose scoring rules are found to either not be proper—as in the case of Good Judgment Open or CSET-Foretell (now INFER)—or incentivize distorting one’s true probabilities to maximize the chances of placing in the top few positions which earn a monetary reward—as in the case of Metaculus. In addition, in almost all cases, forecasting platforms—or, for that matter, prediction markets—disincentivize collaboration.

Against that backdrop, [Reciprocal Scoring: A Method for Forecasting Unanswerable Questions](#), Karger et al. describe a method to elicit predictions in situations in which resolutions are outright not possible, or very far away. They provide some preliminary evidence of its effectiveness in the form of a small randomized trial. However, in the post-peer-review discussion phase in social media, Karger et al.’s method was met with an extremely lukewarm reception from the community of forecasting practitioners, which has grown to view methods which resemble Keynesian Beauty Contests with suspicion.

In this working paper, we outline an alternative incentivization method, an “amplified oracle”, which roughly looks as follows:

- There is a trusted authority which cares about its long-term reputation. This central authority is trusted, but has limited capacity.
- Forecasters are then rewarded according to a scheme where they forecast on all questions, the central authority predicts on only a few questions chosen randomly, and then forecasters are rewarded according to the proximity of their predictions to that central authority.

Towards the middle of the paper, we also point out that the “central authority” doesn’t have to be particularly privileged or more trusted than the rest. The

*Quantified Uncertainty Research Institute.

core feature is rather that some forecasters predict other forecasters and the second group of forecasters predicts reality.

This scheme has the advantage that forecasters can be rewarded speedily for questions which either have no objective resolution or happen far in the future. In addition, if the trusted authority has long horizons, it can be rewarded when the question resolves in the future, or according to the best guess of a future forecasting system.

Description of the method

In the interest of brevity, we shall outline our method by means of an example, and the example shall be the question “Will the People’s Republic of China have annexed at least half of Taiwan by 2050?”, as operationalized by [Metaculus](#). Alas, this method requires a cluster of questions, so the reader should picture a cluster of questions similar to that one.

The oracle determines a rough prior of all questions, in order to reduce potential reward.

Taiwan has been independent of mainland China since the 25th of October 1945, i.e., 76 years into the past. Per Laplace’s law, the chances that this will change by 2050 is $1 - (1 - \frac{1}{(2021-1945)+2})^{2050-2021} \approx 31\%$. Lets take this 31% as the authority’s initial probability. Note that per the [reference class problem](#), other reference classes might have been chosen, so the point of this prior is not to be definitive, but rather to provide a starting point less arbitrary than 50% from which forecaster reward might be computed in the next steps. In the case of a patron aiming to learn from sponsoring a forecasting tournament, the prior might represent the patron’s initial probability.

Forecasters attempt to foresee the oracle’s future forecast.

Forecasters spend some effort trying to come up with forecasts which beat the prior. If they are rewarded in proportion to how much they beat this prior, as outlined in [Beat the house](#), they have an incentive to collaborate.

The oracle predicts on a randomly chosen number of questions

The oracle chooses some questions at random, and produces a forecast for these questions.

Forecasters are rewarded or punished according to how much their probability moves from the prior to the oracle forecast.

As outlined in [Beat the house](#)¹, forecasters end up with a positive balance, if they have moved the probability from the prior towards the oracle’s forecast. But if they moved the probability in the opposite direction, they end up with a negative balance².

In this case, forecasters should be rewarded in proportion to the number of questions which resolve. For instance, if the oracle only looks at one in ten questions, reward or punishment is multiplied by 10. Within a prediction market, the maximum price of a share would be \$10 rather than \$1. But since shares only have a 10% chance of paying out, their price doesn’t exceed \$1.

After a time, the oracle is rewarded or punished

Optionally, after a time, reality is observed, and the oracle itself is paid or punished in proportion to their accuracy. But straight-out rewards require the oracle to have relatively long time horizons, and for anthropic effects to not hold (e.g., predicting a solar flare that ends life has the feature that if you’re wrong, you are already dead).

Alternatively, the oracle can itself be predicting the output of a more thorough investigation. For instance, in the presence of a fair lottery, that investigation can be made very intense and resource intensive with a small probability for relatively little cost.

As a third option, the oracle may instead attempt to predict the client’s forecast after additional time has passed and more information about the matter of interest has surfaced. In this case, it wouldn’t be that the client is a worse forecaster than the oracle, but just busier.

Readers may notice that this proposal resembles previously proposed amplification setups, such as those in [Lagerros et al.](#), which themselves build on ideas by Paul Christiano.

Discussion of the method

More on the central authority

Implicitly, we have been presenting this central authority as more trustworthy, or more truth-seeking than the other forecasters which are used to amplify it. But this might not necessarily be the case at all. This method might instead be used as a cost-saving device. For instance, maybe some forecasters are more willing

¹Though this proposal can also be trivially adapted for use with other scoring rules.

²Or still with a positive balance if the fixed reward is high enough

to receive rewards later rather than immediately, but both sets of forecasters produce similar quality forecasts. Alternatively, perhaps all forecasters have similar discount rates, but rewarding many forecasters for predictions which will be resolved long into the future might be too expensive. In that case, tournament designers might arbitrarily divide a tournament’s cohort into forecasters rewarded now and authorities which are rewarded later, with the former attempting to predict and give information to the latter.

Note that this amplification method might go very wrong³ if forecasting something such as “Will the US dollar suffer from large amounts of inflation?”. In that case, forecasters’ who have reason to believe in inflation might self-select into the group which gets rewarded now (in dollars), and likewise forecasters with beliefs about lack of inflation might self-select into the group which gets rewards later. In that case, the first group of short-termists might sneakily predict that the second group will think that there will be no inflation, but the link between the forecasting system and reality would have been weakened.

We might also want to allow the central authority to see the outputs produced by the forecasters. This allows a forecaster with unique knowledge or exceptional research abilities to predict using all their knowledge, while ensuring that they are not punished for doing so by the relative ignorance of the central authority.

Evidence base and comparison to Karger et al.’s method

Most of the reasons why I think why this method might be superior come from first principles reasoning or from my experience with forecasting. With regard to first-principle reasoning, both steps of the system are incentive compatible, and rational actors would thus be incentivized to predict reality.

With regard to evidence from my experience as a forecaster, this method feels more grounded. A first group of forecasters try to predict or amplify a more patient group of forecasters, and the second group tries to predict reality. If a forecaster has unique information, they can and should try to convince the second group. If the evidence is just very private or hard to communicate, the forecaster with that information could offer a bet to the central authority, and that willingness to bet is itself evidence that the central authority can use. Forecasters are incentivized to conduct expensive searches (e.g., interviews with citizens of far-away lands), cf. [Linch Zhang](#).

In contrast, Karger et al.’s method has weird loops; teams are not aiming to forecast reality, but rather to forecast what the other team will forecast that one’s team will forecast that the other team will forecast... In the presence of Schelling points, human biases, laziness, etc., it is not clear that this process converges to the truth. For instance, the forecaster which puts in the most research effort, or the group which puts in the most effort, is disadvantaged:

³Note that normal forecasting or normal prediction markets will also suffer from this issue if forecasters or market participants use them as hedges.




Figure 1: Proposed design




Figure 2: Karger et al.'s design

ideally, both groups want to put in the same amount of effort, or, equivalently, find out the same things. This is particularly egregious when a forecaster has some piece of evidence that they don't expect anybody on the other team to have. In that case, in Karger et al.'s schema, the forecaster with unique knowledge may want to forecast as if she did not know it.

Some of the authors in Karger et al.'s paper bring forward the argument that if one has large enough teams, each person should expect and equivalent someone in the other team to find the same evidence. Although perhaps true in the limit, this in my experience does not seem likely to be true in any degree of practice.

Karger et al.'s method has the advantage that it produces a legible output, e.g., a wiki. In our case, the first group of forecasters might use similar infrastructure when predicting the forecasts from the second group, and the second group might use that infrastructure to organize its own thoughts. So that doesn't seem like a unique advantage of Karger et al.'s methods.

Conclusion

In conclusion, the proposed method of amplifying an perhaps more expensive and powerful but perhaps just more patient group of forecasters has the benefit of appearing more grounded, and not having the weird loops in Karger et al.'s reciprocal scoring proposal—or in other Keynesian beauty contest designs. We look forward to someone trying to implement it.